

LSTM Stock Prediction Project Plots | ECE 539 | Aidan Pierre-Louis

```
In [1]: import yfinance as yf
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.layers import LSTM, Dense, Dropout
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt

In [2]: np.random.seed(42) # random seed for reproducibility
tf.random.set_seed(42)

# dow jones companies to test model on
dowjones_comps = [
    "AAPL", "MSFT", "AMZN", "GOOGL", # Big Tech # "FB"
    "JNJ", "PFE", "MRK", "UNH", "ABT", # Healthcare
    "PG", "KO", "PEP", "WMT", "HD", # Consumer
    "GS", "JPM", "V", "AXP", "CRM", # Financials
    "DIS", "NFLX", "CMCSA", "T", "INTC", # Media & Telecom
    "NKE", "MCD", "SBUX", "CAT", "MMM", # Industrials
    "BA", "HON", "UNP", "DOW", # Industrials # "UTX"
    "IBM", "CSCO", "AAP", "TRV", "CVX", # Other
    "XOM", "VZ", "MS", "RTX", "WBA" # Other
]

In [3]: # Create lagged features for LSTM
def create_lagged_features(series, n_lags):
    df = pd.DataFrame(series)
    for i in range(1, n_lags + 1):
        df[f'lag_{i}'] = series.shift(i)
    #df.dropna(inplace=True) # Drop rows with NaN values which are the result of shifting, use create_sequences()
    df.fillna(0.0, inplace=True) # replace rows with NaN values from shifting to prevent loss of indices, deprecates
    return df

# Prepare stock data with lagged features
def prepare_stock_data(stock_symbol, n_lags, start='2020-01-01', end='2023-01-01'):
```

```

data = yf.download(stock_symbol, start=start, end=end)
series = data['Close'].dropna()
df = create_lagged_features(series, n_lags)
scaler = MinMaxScaler(feature_range=(-1, 1))
#scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(df)
return scaled_data, scaler, series

# Prepare data, Hyperparam lag features
n_lags = 10
comp_idx = 20 # change the index here for selecting dow jones company to train data on
scaled_data, scaler, series = prepare_stock_data(dowjones_comps[comp_idx], n_lags)

```

```
[*****100%*****] 1 of 1 completed
```

```

In [4]: # Create sequences, deprecated function after altering create_lagged_features to maintain index count
def create_sequences(data, sequence_length):
    X, y = [], []
    for i in range(sequence_length, len(data)):
        X.append(data[i-sequence_length:i, :-1])
        y.append(data[i, -1])
    return np.array(X), np.array(y)

# new function to create sequences
def create_sequences_new(data):
    X, y = data[:, :-1], data[:, -1]
    # print(f'shape of x {len(X)}')
    return np.array(X), np.array(y)

def create_sequences_(data, sequence_length):
    X, y = [], []
    for i in range(len(data)):
        X.append(data[i-sequence_length:i, :-1])
        y.append(data[i, -1])
    return np.array(X), np.array(y)

sequence_length = n_lags
# X, y = create_sequences(scaled_data, sequence_length)
X, y = create_sequences_new(scaled_data)

```

```

In [5]: # Split into training and testing sets
split_percent = 0.9

```

```

#split = int(split_percent * len(X))

#X_train, y_train = X[:split], y[:split]
#X_test, y_test = X[split:], y[split:]

# 70/15/15 split
train_size = int(0.7 * len(X))
val_size = int(0.15 * len(X))
test_size = len(X) - train_size - val_size

X_train, y_train = X[:train_size], y[:train_size]
X_val, y_val = X[train_size:train_size+val_size], y[train_size:train_size+val_size]
X_test, y_test = X[train_size+val_size:], y[train_size+val_size:]

```

```

In [6]: # Reshape data for LSTM
#X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], X_train.shape[2])) # deprecated with create_sequences
#X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], X_test.shape[2])) # deprecated with create_sequences
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1]))
X_val = X_val.reshape((X_val.shape[0], X_val.shape[1]))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1]))

X_train = np.expand_dims(X_train, axis=-1) # Add a new axis for features
X_val = np.expand_dims(X_val, axis=-1) # Add a new axis for features
X_test = np.expand_dims(X_test, axis=-1) # Add a new axis for features

# TODO redundancy, create method to do all this reshaping
first_values = X_train[:, 0]
X_train[:, :, 0] = first_values

first_values = X_val[:, 0]
X_val[:, :, 0] = first_values

first_values = X_test[:, 0]
X_test[:, :, 0] = first_values

```

```

In [7]: #X_test.shape, X_test[:10, :, :]

```

```

In [8]: # Hyperparam, dropout prob
p = 0.2
hidden_units = 50

```

```
# Build the LSTM model
model = Sequential([
    LSTM(hidden_units, return_sequences=True, input_shape=(X_train.shape[1], X_train.shape[2])),
    Dropout(p),
    LSTM(hidden_units, return_sequences=False),
    Dropout(p),
    Dense(25),
    Dense(1)
])
```


C:\Users\p_pie\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\keras\src\layers\rnn\rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```
super().__init__(**kwargs)
```


```
In [9]: # Hyperparameters
learning_rate = 0.001
batch_size = 16
epochs = 100


optimizer = Adam(learning_rate=learning_rate)
# Compile the model
model.compile(optimizer=optimizer, loss='mean_squared_error', metrics=['accuracy'])
# Train the model
history = model.fit(X_train, y_train, batch_size=batch_size, epochs=epochs, validation_data=(X_val, y_val), verbose=)
```


```
Epoch 1/100
34/34 ————— 2s 12ms/step - accuracy: 0.0011 - loss: 0.0990 - val_accuracy: 0.0000e+00 - val_loss: 0.0197
Epoch 2/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0271 - val_accuracy: 0.0000e+00 - val_loss: 0.0482
Epoch 3/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0266 - val_accuracy: 0.0000e+00 - val_loss: 0.1123
Epoch 4/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 0.1298
Epoch 5/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0271 - val_accuracy: 0.0000e+00 - val_loss: 0.2623
Epoch 6/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0276 - val_accuracy: 0.0000e+00 - val_loss: 0.2932
Epoch 7/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0287 - val_accuracy: 0.0000e+00 - val_loss: 0.4000
Epoch 8/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0299 - val_accuracy: 0.0000e+00 - val_loss: 0.2615
Epoch 9/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0271 - val_accuracy: 0.0000e+00 - val_loss: 0.2776
Epoch 10/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0288 - val_accuracy: 0.0000e+00 - val_loss: 0.4865
Epoch 11/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0281 - val_accuracy: 0.0000e+00 - val_loss: 0.4892
Epoch 12/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0277 - val_accuracy: 0.0000e+00 - val_loss: 0.4251
Epoch 13/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 0.4612
Epoch 14/100
34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0289 - val_accuracy: 0.0000e+00 - val_loss: 0.3807
```


Epoch 15/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0282 - val_accuracy: 0.0000e+00 - val_loss: 0.3607


Epoch 16/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0279 - val_accuracy: 0.0000e+00 - val_loss: 0.4600


Epoch 17/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0275 - val_accuracy: 0.0000e+00 - val_loss: 0.6637


Epoch 18/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0278 - val_accuracy: 0.0000e+00 - val_loss: 0.4633


Epoch 19/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0266 - val_accuracy: 0.0000e+00 - val_loss: 0.5935


Epoch 20/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0275 - val_accuracy: 0.0000e+00 - val_loss: 0.7084


Epoch 21/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0269 - val_accuracy: 0.0000e+00 - val_loss: 0.5997


Epoch 22/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0279 - val_accuracy: 0.0000e+00 - val_loss: 0.6847


Epoch 23/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0282 - val_accuracy: 0.0000e+00 - val_loss: 0.7260


Epoch 24/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0274 - val_accuracy: 0.0000e+00 - val_loss: 0.9658


Epoch 25/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0296 - val_accuracy: 0.0000e+00 - val_loss: 0.6441


Epoch 26/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0269 - val_accuracy: 0.0000e+00 - val_loss: 0.6576


Epoch 27/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0269 - val_accuracy: 0.0000e+00 - val_loss: 0.4974


Epoch 28/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0271 - val_accuracy: 0.0000e+00 - val_loss: 0.6707


Epoch 29/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0269 - val_accuracy: 0.0000e+00 - val_loss: 0.8756


Epoch 30/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0291 - val_accuracy: 0.0000e+00 - val_loss: 0.7826


Epoch 31/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 0.7941


Epoch 32/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0289 - val_accuracy: 0.0000e+00 - val_loss: 0.6045


Epoch 33/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0275 - val_accuracy: 0.0000e+00 - val_loss: 0.8144


Epoch 34/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0271 - val_accuracy: 0.0000e+00 - val_loss: 0.7256


Epoch 35/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0275 - val_accuracy: 0.0000e+00 - val_loss: 0.8802


Epoch 36/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0287 - val_accuracy: 0.0000e+00 - val_loss: 0.5748


Epoch 37/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0255 - val_accuracy: 0.0000e+00 - val_loss: 0.6891


Epoch 38/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0284 - val_accuracy: 0.0000e+00 - val_loss: 0.9506


Epoch 39/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0287 - val_accuracy: 0.0000e+00 - val_loss: 0.7354


Epoch 40/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0278 - val_accuracy: 0.0000e+00 - val_loss: 0.7279


Epoch 41/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0269 - val_accuracy: 0.0000e+00 - val_loss: 0.8300


Epoch 42/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0294 - val_accuracy: 0.0000e+00 - val_loss: 1.1028


Epoch 43/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0289 - val_accuracy: 0.0000e+00 - val_loss: 1.0312


Epoch 44/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 0.8961


Epoch 45/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0279 - val_accuracy: 0.0000e+00 - val_loss: 0.9732


Epoch 46/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0268 - val_accuracy: 0.0000e+00 - val_loss: 0.8243


Epoch 47/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0265 - val_accuracy: 0.0000e+00 - val_loss: 0.9001


Epoch 48/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0267 - val_accuracy: 0.0000e+00 - val_loss: 1.1332


Epoch 49/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0286 - val_accuracy: 0.0000e+00 - val_loss: 1.0011


Epoch 50/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0273 - val_accuracy: 0.0000e+00 - val_loss: 1.2692


Epoch 51/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0276 - val_accuracy: 0.0000e+00 - val_loss: 0.9476


Epoch 52/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0287 - val_accuracy: 0.0000e+00 - val_loss: 1.0026


Epoch 53/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0269 - val_accuracy: 0.0000e+00 - val_loss: 1.1197


Epoch 54/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0265 - val_accuracy: 0.0000e+00 - val_loss: 0.9661


Epoch 55/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0270 - val_accuracy: 0.0000e+00 - val_loss: 1.0545


Epoch 56/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 1.4728


Epoch 57/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0296 - val_accuracy: 0.0000e+00 - val_loss: 0.7464


Epoch 58/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0271 - val_accuracy: 0.0000e+00 - val_loss: 1.1494


Epoch 59/100
34/34  0s 4ms/step - accuracy: 0.0011 - loss: 0.0273 - val_accuracy: 0.0000e+00 - val_loss: 1.0920


Epoch 60/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0258 - val_accuracy: 0.0000e+00 - val_loss: 1.3160


Epoch 61/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0266 - val_accuracy: 0.0000e+00 - val_loss: 0.9402


Epoch 62/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0279 - val_accuracy: 0.0000e+00 - val_loss: 0.9763


Epoch 63/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0273 - val_accuracy: 0.0000e+00 - val_loss: 1.2052


Epoch 64/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0277 - val_accuracy: 0.0000e+00 - val_loss: 0.7822


Epoch 65/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0282 - val_accuracy: 0.0000e+00 - val_loss: 1.3277


Epoch 66/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0294 - val_accuracy: 0.0000e+00 - val_loss: 1.1263


Epoch 67/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0279 - val_accuracy: 0.0000e+00 - val_loss: 1.3186


Epoch 68/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 1.2020


Epoch 69/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0285 - val_accuracy: 0.0000e+00 - val_loss: 1.1563


Epoch 70/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0287 - val_accuracy: 0.0000e+00 - val_loss: 1.5343


Epoch 71/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0276 - val_accuracy: 0.0000e+00 - val_loss: 1.2984


Epoch 72/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0262 - val_accuracy: 0.0000e+00 - val_loss: 1.5258


Epoch 73/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0277 - val_accuracy: 0.0000e+00 - val_loss: 1.4791


Epoch 74/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0279 - val_accuracy: 0.0000e+00 - val_loss: 1.4641


Epoch 75/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0272 - val_accuracy: 0.0000e+00 - val_loss: 1.2697


Epoch 76/100
34/34  0s 4ms/step - accuracy: 0.0011 - loss: 0.0254 - val_accuracy: 0.0000e+00 - val_loss: 1.3089


Epoch 77/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0265 - val_accuracy: 0.0000e+00 - val_loss: 1.2898


Epoch 78/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 1.2186


Epoch 79/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 1.4723


Epoch 80/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0278 - val_accuracy: 0.0000e+00 - val_loss: 1.3823


Epoch 81/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0276 - val_accuracy: 0.0000e+00 - val_loss: 1.5285


Epoch 82/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0262 - val_accuracy: 0.0000e+00 - val_loss: 1.2899


Epoch 83/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0271 - val_accuracy: 0.0000e+00 - val_loss: 1.4331


Epoch 84/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0259 - val_accuracy: 0.0000e+00 - val_loss: 1.2393


Epoch 85/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0269 - val_accuracy: 0.0000e+00 - val_loss: 1.0645


Epoch 86/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0267 - val_accuracy: 0.0000e+00 - val_loss: 0.8009


Epoch 87/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0258 - val_accuracy: 0.0000e+00 - val_loss: 0.9577


Epoch 88/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0265 - val_accuracy: 0.0000e+00 - val_loss: 1.2779


Epoch 89/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0270 - val_accuracy: 0.0000e+00 - val_loss: 1.1045


Epoch 90/100
34/34  0s 4ms/step - accuracy: 0.0011 - loss: 0.0270 - val_accuracy: 0.0000e+00 - val_loss: 1.1121


Epoch 91/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0287 - val_accuracy: 0.0000e+00 - val_loss: 0.9110


Epoch 92/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0262 - val_accuracy: 0.0000e+00 - val_loss: 1.1398


Epoch 93/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0274 - val_accuracy: 0.0000e+00 - val_loss: 1.1613

Epoch 94/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0265 - val_accuracy: 0.0000e+00 - val_loss: 0.9715

Epoch 95/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0280 - val_accuracy: 0.0000e+00 - val_loss: 0.9373

Epoch 96/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0274 - val_accuracy: 0.0000e+00 - val_loss: 0.8317

Epoch 97/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0278 - val_accuracy: 0.0000e+00 - val_loss: 0.8443

Epoch 98/100
34/34  0s 3ms/step - accuracy: 0.0011 - loss: 0.0267 - val_accuracy: 0.0000e+00 - val_loss: 0.9358

Epoch 99/100

34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0275 - val_accuracy: 0.0000e+00 - val_loss: 1.0500

Epoch 100/100

34/34 ————— 0s 3ms/step - accuracy: 0.0011 - loss: 0.0278 - val_accuracy: 0.0000e+00 - val_loss: 0.7706

```
In [10]: # Make predictions
predictions = model.predict(X_test)
predictions = scaler.inverse_transform(np.concatenate((np.zeros((predictions.shape[0], n_lags)), predictions), axis=1))
```

4/4 ————— 0s 42ms/step

```
In [11]: # Forecast next forecast_length days using the days from the test set, Hyperparam forecast length
forecast_length = 30
input_seq = X_test[-1]






















# function deprecated with create_sequences
def forecast_predictions(forecast_length=30, input_seq=[]):
    predictions = []
    for _ in range(forecast_length):
        print(f'Forecast step {_+1}:')
        pred = model.predict(input_seq.reshape(1, sequence_length, X_test.shape[2]))
        new_row = np.append(input_seq[1:, -1], pred) # Append the new prediction
        input_seq = np.vstack((input_seq[1:], new_row)) # Update the input sequence with the new prediction
        pred_inverse_scaled = scaler.inverse_transform(np.concatenate((np.zeros((1, n_lags)), pred), axis=1))[:, -1]
        predictions.append(pred_inverse_scaled[0])
    return predictions

# forecasted_values = forecast_predictions(forecast_length, input_seq)

def forecast_predictions_new(model, data, sequence_length, forecast_length, scaler):
    predictions = []
    last_sequence = data[-n_lags:] # Take the last n_lags closing prices as the initial input
    for _ in range(forecast_length):
        print(f'Forecast step {_+1}:')
        # Reshape last_sequence to match the model input shape (1, sequence_length, 1)
        last_sequence = last_sequence.reshape((1, sequence_length, 1))
        # Predict the next value
        next_value = model.predict(last_sequence)[0, 0]
```

```
# Append the prediction
predictions.append(next_value)
# Update the sequence, rolling it forward
next_value = np.array([[next_value]]) # Reshape next_value to be 3D: (1, 1, 1)
last_sequence = np.append(last_sequence[:, 1:, :], next_value.reshape(1, 1, 1), axis=1)
# Inverse scale the predictions after the loop
predictions = np.array(predictions).reshape(-1, 1)
future_predictions = scaler.inverse_transform(np.concatenate((np.zeros((len(predictions), sequence_length)), predictions), axis=1))
return future_predictions

# Use the model to forecast the next n_future days
forecasted_values = forecast_predictions_new(model, input_seq, sequence_length, forecast_length, scaler)
```

Forecast step 1:
1/1  0s 13ms/step
Forecast step 2:
1/1  0s 11ms/step
Forecast step 3:
1/1  0s 10ms/step
Forecast step 4:
1/1  0s 10ms/step
Forecast step 5:
1/1  0s 11ms/step
Forecast step 6:
1/1  0s 10ms/step
Forecast step 7:
1/1  0s 10ms/step
Forecast step 8:
1/1  0s 11ms/step
Forecast step 9:
1/1  0s 11ms/step
Forecast step 10:
1/1  0s 10ms/step
Forecast step 11:
1/1  0s 10ms/step
Forecast step 12:
1/1  0s 10ms/step
Forecast step 13:
1/1  0s 10ms/step
Forecast step 14:
1/1  0s 11ms/step
Forecast step 15:
1/1  0s 12ms/step
Forecast step 16:
1/1  0s 10ms/step
Forecast step 17:
1/1  0s 11ms/step
Forecast step 18:
1/1  0s 11ms/step
Forecast step 19:
1/1  0s 10ms/step
Forecast step 20:
1/1  0s 10ms/step
Forecast step 21:
1/1  0s 11ms/step

```

Forecast step 22:
1/1 _____ 0s 10ms/step
Forecast step 23:
1/1 _____ 0s 15ms/step
Forecast step 24:
1/1 _____ 0s 12ms/step
Forecast step 25:
1/1 _____ 0s 11ms/step
Forecast step 26:
1/1 _____ 0s 11ms/step
Forecast step 27:
1/1 _____ 0s 11ms/step
Forecast step 28:
1/1 _____ 0s 12ms/step
Forecast step 29:
1/1 _____ 0s 11ms/step
Forecast step 30:
1/1 _____ 0s 11ms/step

```

```

In [12]: # Split the series into training, validation, and test parts
train = series[:train_size]
val = series[train_size:train_size+val_size]
test = series[train_size+val_size:]

# if slicing logic is wrong due to integer math
if len(val) != len(test):
    train = series[:train_size-1]
    val = series[train_size:train_size+val_size+1]
    test = series[train_size+val_size:]

# Convert to DataFrame
train = train.to_frame()
val = val.to_frame()
test = test.to_frame()

# Ensure `valid` is renamed to `val`
# Generate predictions for validation and forecast as needed
# Add predictions to the validation set
test['Predictions'] = np.nan
test['Predictions'].iloc[:len(predictions)] = predictions

# Create forecast dates (adjust if needed)

```

```

forecast_dates = pd.date_range(start=test.index[-1], periods=forecast_length + 1, inclusive='right')

# Create a DataFrame for forecasted values
forecast = pd.DataFrame(forecasted_values, index=forecast_dates, columns=['Predictions'])

# Plot the results
plt.figure(figsize=(16, 8))
plt.title(f'Stock Price Prediction {dowjones_comps[comp_idx]}')
plt.xlabel('Date')
plt.ylabel('Close Price USD ($)')
plt.plot(train, label='Train')
plt.plot(val, label='Validation')
plt.plot(test['Close'], label='Test Data', linestyle='-', color='b')
plt.plot(test['Predictions'], label='Test Predictions', linestyle='--', color='r')
plt.plot(forecast, label='Forecast', linestyle='--', color='g')
plt.legend(loc='lower right')
plt.show()

```

C:\Users\p_pie\AppData\Local\Temp\ipykernel_18844\610895317.py:21: FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!


You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) this will never work to update the original DataFrame or Series, because the intermediate object on which we are setting values will behave as a copy. A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```


Use `df.loc[row_indexer, "col"] = values` instead, to perform the assignment in a single step and ensure this keeps updating the original `df`.

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
test['Predictions'].iloc[:len(predictions)] = predictions
```

```
In [13]: # Plot the results
plt.figure(figsize=(16, 8))
plt.title(f'Stock Price Prediction {dowjones_comps[comp_idx]}')
plt.xlabel('Date')
plt.ylabel('Close Price USD ($)')
plt.plot(val, label='Validation')
plt.plot(test['Close'], label='Test Data', linestyle='--', color='b')
plt.plot(test['Predictions'], label='Test Predictions', linestyle='--', color='r')
plt.plot(forecast, label='Forecast', linestyle='--', color='g')
plt.legend(loc='lower right')
plt.show()
```




```
In [14]: # Extract loss and accuracy
loss = history.history['loss']
val_loss = history.history['val_loss']
accuracy = history.history.get('accuracy', []) # If your model records accuracy
val_accuracy = history.history.get('val_accuracy', [])

# Compute test loss
results = model.evaluate(X_val, y_val, batch_size=batch_size, verbose=0)
test_loss = results[0]

# Plot Training and Validation Loss
plt.figure(figsize=(12, 6))
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
```

```
plt.axhline(y=test_loss, color='r', linestyle='--', label=f'Test Loss ({test_loss:.4f})')
plt.title(f'Training and Validation Loss {dowjones_comps[comp_idx]}')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [15]: # Plot Training and Validation Accuracy (if applicable)
if accuracy and val_accuracy:
    plt.figure(figsize=(12, 6))
    plt.plot(accuracy, label='Training Accuracy')
    plt.plot(val_accuracy, label='Validation Accuracy')
    plt.title(f'Training and Validation Accuracy {dowjones_comps[comp_idx]}')
    plt.xlabel('Epochs')
```

```
plt.ylabel('Accuracy')  
plt.legend()  
plt.show()
```

